

Core Architecture & Data Model

| Feature | Redis Enterprise | GridGain Enterprise |
|--------------------------|--|---|
| Architecture | Shared-nothing, in-memory, sharded | Distributed MPP (Massively Parallel Processing) architecture |
| Data model | Key-value, data structures (List, Set, Hash, JSON, Stream, TimeSeries) | Key-value + full ANSI SQL tables + ACID transactions + colocated compute |
| Memory model | DRAM-based with optional SSD tiering | Off-heap memory with durable memory (RAM+Disk) |
| Data ownership | Cache-first (copy of truth); apps control sync | In-memory data grid (primary store of record) |
| Data partitioning | Hash-based sharding | Affinity-based partitioning with compute/data co-location |

Transactions & Consistency

| Feature | Redis Enterprise | GridGain Enterprise |
|-------------------------------|---|---|
| Transaction support | Single-shard, optimistic multi-command transactions | Distributed, ACID-compliant transactions with 2PC (two-phase commit) |
| Locking model | None (optimistic locking only) | Pessimistic & optimistic locking, deadlock prevention |
| Cross-node consistency | Not supported | Fully supported with global coordination |
| JTA/XA support | Not supported | Supported |

Persistence & Durability

| Feature | Redis Enterprise | GridGain Enterprise |
|----------------------------|---------------------------------|--|
| Persistence options | RDB snapshots + AOF logs | Native persistence , write-ahead logging, checkpointing |
| Startup behavior | Requires reload from disk | Instant-on recovery from durable memory |
| Data loss on crash | Possible unless carefully tuned | Zero data loss with journaling and checkpointing |

Querying & Analytics

| Feature | Redis Enterprise | GridGain Enterprise |
|-------------------------|--|--|
| Query language | Redis commands, custom APIs (Search/JSON/TimeSeries) | ANSI SQL (DDL/DML), JDBC/ODBC , full relational support |
| Secondary indexes | Yes, via RedisSearch | Yes, including distributed indexes |
| Join support | No | Yes, distributed joins supported |
| Aggregation & filtering | Yes (limited) | Yes (SQL GROUP BY, HAVING, etc.) |
| Analytics engine | Limited | Distributed MPP compute |

Compute & Streaming

| Feature | Redis Enterprise | GridGain Enterprise |
|---------------------|---------------------------------------|--|
| Pub/Sub support | Yes | Yes |
| Streaming engine | Redis Streams | Built-in stream processor + Kafka/Spark integration |
| Server-side compute | No (uses Lua scripts or client logic) | Yes, collocated processing (map/reduce, ML, continuous queries) |
| ML integration | External models only | Integrated with Apache Ignite ML + collocated scoring |

Clustering & Scalability

| Feature | Redis Enterprise | GridGain Enterprise |
|-----------------|-----------------------------|---------------------------------------|
| Cluster type | Shared-nothing, shard-based | Peer-to-peer, partitioned MPP cluster |
| Sharding | Auto-sharding, hash-based | Affinity-aware partitioning |
| Replication | Active-active with CRDTs | Synchronous/asynchronous replication |
| Elastic scaling | Supported | Supported with automatic rebalancing |

| Feature | Redis Enterprise | GridGain Enterprise |
|--------------------------|------------------|--|
| Compute/data co-location | Not supported | Fully supported for performance gains |

Security, Monitoring, Tooling

| Feature | Redis Enterprise | GridGain Enterprise |
|-------------------------|---|--|
| Security | TLS, RBAC, auditing, compliance (SOC2, HIPAA) | TLS, RBAC, auditing, LDAP, Kerberos, fine-grained access control |
| Monitoring & management | Redis Insight, CLI, REST | GridGain Control Center for monitoring, config, snapshot management |
| Backup & restore | Manual/automated snapshotting | Point-in-time recovery , snapshot scheduling |
| Deployment options | Cloud, Kubernetes, hybrid | On-prem, hybrid, Kubernetes, cloud-native |

Use Cases (Investment Banking & Markets)

| Application Area | Redis Enterprise | GridGain Enterprise |
|----------------------------|--|--|
| Market data cache | Sub-ms pricing cache, tick data lookup | Used as upstream cache, but better suited for risk calcs |
| Trade matching / lifecycle | Not ACID-safe for cross-shard trade ops | Yes – full ledger handling with ACID |
| Fraud detection | Real-time feature store + ML model scoring | Inline model execution with stream + collocated compute |
| Risk dashboards | Limited querying (filter + aggregate only) | Distributed SQL-based aggregation and analytics |
| Settlement & clearing | Not designed for transactional guarantees | Yes – fully transactional updates, snapshot, failover |

Final Recommendation Summary

| Category | Best Choice |
|---------------------------|------------------|
| Speed for key-value cache | Redis Enterprise |

| Category | Best Choice |
|--------------------------------|---------------------|
| Distributed ACID transactions | GridGain Enterprise |
| Real-time dashboards + SQL | GridGain Enterprise |
| ML feature serving | Redis Enterprise |
| Streaming analytics + compute | GridGain Enterprise |
| Ease of use / Dev onboarding | Redis Enterprise |
| Full data platform replacement | GridGain Enterprise |

Redis Enterprise –

Not Suitable Use Cases Compared to GridGain

| Use Case / Scenario | Why Redis Falls Short | Why GridGain Excels |
|---|---|--|
| 1. Distributed ACID Transactions | Redis only supports single-shard transactions with no cross-node ACID compliance. | GridGain supports distributed ACID transactions across nodes via two-phase commit. Ideal for banking, trade matching, ledger ops. |
| 2. Complex SQL-Based Analytics | Redis lacks native SQL engine; RedisSearch is limited for joins, group-bys, and relational logic. | GridGain supports full ANSI SQL , joins, aggregation, subqueries, and JDBC/ODBC access. Used for VaR, intraday P&L, exposure reporting. |
| 3. Co-located Compute and Data Processing | Redis has no concept of co-located server-side compute; analytics must happen outside. | GridGain enables in-grid computation (map-reduce, ML scoring, streaming) on the same nodes as data — reducing network latency. |
| 4. Trade Lifecycle Management (IBOR) | Redis can't guarantee atomic multi-key/multi-entity updates across shards. | GridGain enables transactional ledger updates , with snapshotting and recovery. Ideal for matching, booking, settlement pipelines. |
| 5. T+0 Settlement & Streaming Event Processing | Redis Streams lacks integration with Spark/Kafka for true streaming ETL. | GridGain includes native streaming support , windowed aggregation, and integration with Kafka/Spark for real-time clearing ops. |
| 6. Real-Time Risk Aggregation | No SQL-based aggregations across datasets; slow or manual workaround required. | GridGain's SQL engine + MPP architecture enables real-time, in-memory VaR, stress tests, and exposure tracking . |

| Use Case / Scenario | Why Redis Falls Short | Why GridGain Excels |
|---|--|--|
| 7. Distributed Ledger with Durability | Redis persistence is snapshot/AOF; requires reloads; risk of partial writes. | GridGain uses native persistence + journaling + checkpoints for instant recovery and durability , even across thousands of trades. |
| 8. Enterprise Integration (DIH) | Redis requires external ETL for RDBMS/NoSQL integration. | GridGain acts as a Digital Integration Hub , ingesting from Oracle, PostgreSQL, Kafka, Hive, etc., and serving data via SQL. |
| 9. Complex Compliance Auditing or Governance | Redis lacks snapshot-level auditability and recovery tooling. | GridGain has built-in snapshotting, backup, audit logs , and RBAC integrated into the Control Center. |
| 10. Highly Structured, High-Volume Data Models | Redis performs poorly with complex, relational models (e.g., trades, allocations, counterparties). | GridGain allows schema-rich models , joins, and storage of high-volume structured financial data (like nested trade trees or books). |

Example: Redis Limitations in Financial Systems

| Application | Redis Limitation | Real Consequence |
|----------------------------------|---|---|
| Trade Matching Engine | Cannot guarantee atomic update of order book and ledger across shards | Risk of mismatched trades or partial booking |
| Intraday VaR Engine | Cannot perform distributed joins or SQL aggregations | Incomplete or delayed risk reporting |
| Cross-Account Settlement | No multi-entity atomicity or durability guarantees | Funds debited but not credited—violating regulatory rules |
| Order Audit Trail | No native snapshots or timestamped data history | Difficult to comply with MiFID II, SEC Rule 17a-4 |
| Real-Time Fraud Detection | Features must be pre-joined or precomputed externally | Added latency, inconsistent scoring |

When to Use GridGain Over Redis

Choose **GridGain Enterprise** instead of **Redis Enterprise** when:

- You need **distributed transactions and SQL** on **structured datasets**.
- You're building a **financial ledger, risk engine, or trading system**.
- Your architecture must support **in-memory persistence, streaming analytics, and instant recovery**.
- You're integrating **multiple data sources** (e.g. RDBMS + Kafka + APIs).
- You want a **unified platform** for compute + storage + ML pipelines.

Redis Enterprise –

Best Fit Compared to GridGain

| Use Case / Scenario | Why Redis Excels | Why GridGain Falls Short |
|--|---|---|
| 1. Ultra-Low Latency Caching | Redis delivers sub-millisecond reads/writes using in-memory storage + optimized modules. | GridGain is fast but adds overhead due to distributed processing layers. |
| 2. Simple Key-Value or Session Storage | Ideal for session state, JWT tokens, shopping carts, rate limiting , etc. | GridGain is too heavyweight and complex for these small-scale tasks. |
| 3. Microservice Response Acceleration | Used as a shared cache for microservices (e.g. token auth, config lookups). | GridGain introduces latency and ops complexity for simple lookups. |
| 4. Real-Time Leaderboards / Counters | Native support for atomic INCR, ZADD, ZSCORE, etc. , perfect for gaming or analytics counters. | GridGain lacks native data types like sorted sets; requires more design effort. |
| 5. Pub/Sub Messaging / Streams | Redis Streams and Pub/Sub support lightweight messaging across services with low overhead. | GridGain is designed for streaming + compute, not simple fanout. |
| 6. Real-Time Feature Store for ML | RedisJSON, RedisVector, RedisSearch modules allow fast feature retrieval for model serving. | GridGain can host ML logic but lacks specialized modules for vector search or JSON parsing. |
| 7. Temporary Caching Layer for Expensive Queries | Acts as a read-through/write-through cache over RDBMS or APIs. | GridGain can do this but takes more setup, especially for short-lived data. |
| 8. Time-Series Analytics (Basic) | RedisTimeSeries module supports rollups, downsampling, compact storage . | GridGain supports streaming but is heavier to manage for basic time-series. |
| 9. API Rate Limiting / Throttling | Native support for token bucket & sliding window logic in-memory. | GridGain would require external logic + cluster overhead. |
| 10. SaaS Tenant Isolation | Namespace isolation via logical databases; can scale per tenant easily. | GridGain is harder to multitenant at small scale. |

GridGain Overkill or Not a Fit For:

| Application | Why Redis Is Better | GridGain Drawback |
|---------------------------------------|---|--|
| Config cache for microservices | Redis has built-in TTLs, instant lookup, and high availability. | GridGain requires cluster coordination, setup for trivial data. |
| Leaderboards in gaming/ads | Redis Sorted Sets = native fit; optimized for atomic ranking ops. | GridGain lacks native sorted-set functionality. |
| IoT device telemetry lookup | Redis delivers <1ms lookups for millions of devices. | GridGain is better suited for complex event processing, not key reads. |
| SaaS user preference storage | Redis works well for per-user metadata, profiles, sessions. | GridGain is excessive for 1–2 KB profile blobs. |
| Real-time fraud scoring API | RedisFeatureStore serves vector/JSON features in <1ms. | GridGain needs full model/feature ingestion or external pipeline. |

When to Use Redis Over GridGain

Choose **Redis Enterprise** instead of **GridGain Enterprise** when:

- You need **low-latency data access** (sub-millisecond) for high-throughput APIs.
- Your workload is **read-heavy** or focused on **caching, session storage, or real-time scoring**.
- You're building **stateless, microservice-based systems**.
- You don't need **SQL, distributed transactions, or streaming analytics pipelines**.
- You want **easy deployment**, lower ops cost, and rapid development cycles.

Summary: When Redis Is the Better Fit

| Redis Enterprise Best For | GridGain Overkill Because... |
|--|--|
| Microservices, token caches, and auth state | Requires no SQL or ACID |
| Leaderboards, real-time counters | Native commands like INCR, ZADD, ZREVRANK |
| Pub/Sub or fanout messaging between components | GridGain's compute grid is too complex for simple messages |
| Feature stores for real-time inference | RedisVector and JSON modules offer instant lookup |

Redis Enterprise Best For

Rate limiting, quotas, API TTL
enforcement

Time-series with rollups and retention

GridGain Overkill Because...

Built-in rate-limiting algorithms

Simple and fast with RedisTimeSeries